

TP2_MLIP

January 21, 2020

```
[1]: %load_ext autoreload
      %autoreload 2

      %matplotlib inline
```

0.0.1 Curso-Taller 2019: Aprendizaje Automático e Imágenes en Python

Trabajo Practico 2 Integrantes: * Ferreyra, Mario Ezequiel (ferreyramario7@gmail.com)
* Gonzalez, Maria Laura (m.laugonzalez@gmail.com) * Kopp, Matias Nicolas (koppmatias97@gmail.com)

```
[2]: import numpy as np
      import matplotlib
      import matplotlib.pyplot as plt
      import seaborn as sns
      import tensorflow as tf
      import keras

      from PIL import Image
      from sklearn.preprocessing import LabelEncoder, OneHotEncoder
      from keras.layers import Input, Dense, Conv2D, MaxPooling2D, Flatten,
      ↪BatchNormalization, Activation, Dropout
      from keras.models import Model, Sequential
      from keras.callbacks import EarlyStopping, ModelCheckpoint, Callback
      from keras.optimizers import RMSprop
      from keras.preprocessing.image import ImageDataGenerator
      from tensorflow.examples.tutorials.mnist import input_data

      from IPython.display import Image as IPImage

      import warnings
      warnings.filterwarnings("ignore")
```

Using TensorFlow backend.

```
[3]: print(tf.__version__)
      print(keras.__version__)
```

1.15.0
2.2.5

```
[4]: sns.set_style("white")
```

Ejercicio 1

Cargamos los datos del MNIST

```
[5]: # importo y guardo MNIST data
mnist = input_data.read_data_sets("./mnist/data/", one_hot=True)
```

WARNING:tensorflow:From <ipython-input-5-5933df94991b>:2: read_data_sets (from tensorflow.contrib.learn.python.learn.datasets.mnist) is deprecated and will be removed in a future version.

Instructions for updating:

Please use alternatives such as official/mnist/dataset.py from tensorflow/models.

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/tensorflow_core/contrib/learn/python/learn/datasets/mnist.py:260: maybe_download (from tensorflow.contrib.learn.python.learn.datasets.base) is deprecated and will be removed in a future version.

Instructions for updating:

Please write your own downloading logic.

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/tensorflow_core/contrib/learn/python/learn/datasets/mnist.py:262: extract_images (from tensorflow.contrib.learn.python.learn.datasets.mnist) is deprecated and will be removed in a future version.

Instructions for updating:

Please use tf.data to implement this functionality.

Extracting ./mnist/data/train-images-idx3-ubyte.gz

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/tensorflow_core/contrib/learn/python/learn/datasets/mnist.py:267: extract_labels (from tensorflow.contrib.learn.python.learn.datasets.mnist) is deprecated and will be removed in a future version.

Instructions for updating:

Please use tf.data to implement this functionality.

Extracting ./mnist/data/train-labels-idx1-ubyte.gz

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/tensorflow_core/contrib/learn/python/learn/datasets/mnist.py:110: dense_to_one_hot (from tensorflow.contrib.learn.python.learn.datasets.mnist) is deprecated and will be removed in a future version.

Instructions for updating:

Please use tf.one_hot on tensors.

Extracting ./mnist/data/t10k-images-idx3-ubyte.gz

Extracting ./mnist/data/t10k-labels-idx1-ubyte.gz

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-

packages/tensorflow_core/contrib/learn/python/learn/datasets/mnist.py:290:
DataSet.__init__ (from tensorflow.contrib.learn.python.learn.datasets.mnist) is
deprecated and will be removed in a future version.
Instructions for updating:
Please use alternatives such as official/mnist/dataset.py from
tensorflow/models.

```
[6]: X_train = mnist.train.images
      y_train = mnist.train.labels

      X_val = mnist.validation.images
      y_val = mnist.validation.labels

      X_test = mnist.test.images
      y_test = mnist.test.labels
```

```
[7]: print(f"X_train shape: {X_train.shape}")
      print(f"y_train shape: {y_train.shape}")
      print()
      print(f"X_val shape : {X_val.shape}")
      print(f"y_val shape : {y_val.shape}")
      print()
      print(f"X_test shape : {X_test.shape}")
      print(f"y_test shape : {y_test.shape}")
```

```
X_train shape: (55000, 784)
y_train shape: (55000, 10)
```

```
X_val shape : (5000, 784)
y_val shape : (5000, 10)
```

```
X_test shape : (10000, 784)
y_test shape : (10000, 10)
```

Ejercicio2

Defina dos redes totalmente conectadas diferentes que sean capaces de clasificar la base MNIST.

Una con 3 capas ocultas y la otra con 5.

Evalúe posible overfitting en ambas y compárelas usándola parte de entrenamiento y la parte de la base destinada para validación.

```
[ ]:
```

Red Fully Connected - 3 Capas Ocultas

```
[8]: model_3hl = Sequential()
model_3hl.add(Dense(512, input_shape=(784,), activation="relu"))
model_3hl.add(Dense(1024, activation="relu"))
model_3hl.add(Dense(512, activation="relu"))
model_3hl.add(Dense(10, activation="softmax"))

model_3hl.compile(
    loss='categorical_crossentropy',
    metrics=['accuracy'],
    optimizer='adam'
)
```

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:66: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:541: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:4432: The name tf.random_uniform is deprecated. Please use tf.random.uniform instead.

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/keras/optimizers.py:793: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:3576: The name tf.log is deprecated. Please use tf.math.log instead.

```
[9]: BATCH_SIZE = 128
NUM_EPOCHS = 20

history_3hl = model_3hl.fit(
    X_train, y_train,
    batch_size=BATCH_SIZE,
    epochs=NUM_EPOCHS,
    verbose=1,
    validation_data=(X_val, y_val)
)
```

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/tensorflow_core/python/ops/math_grad.py:1424: where (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use `tf.where` in 2.0, which has the same broadcast rule as `np.where`

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:1033: The name `tf.assign_add` is deprecated. Please use `tf.compat.v1.assign_add` instead.

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:1020: The name `tf.assign` is deprecated. Please use `tf.compat.v1.assign` instead.

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:3005: The name `tf.Session` is deprecated. Please use `tf.compat.v1.Session` instead.

Train on 55000 samples, validate on 5000 samples

Epoch 1/20

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:190: The name `tf.get_default_session` is deprecated. Please use `tf.compat.v1.get_default_session` instead.

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:197: The name `tf.ConfigProto` is deprecated. Please use `tf.compat.v1.ConfigProto` instead.

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:207: The name `tf.global_variables` is deprecated. Please use `tf.compat.v1.global_variables` instead.

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:216: The name `tf.is_variable_initialized` is deprecated. Please use `tf.compat.v1.is_variable_initialized` instead.

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:223: The name `tf.variables_initializer` is deprecated. Please use `tf.compat.v1.variables_initializer` instead.

55000/55000 [=====] - 10s 174us/step - loss: 0.2100 - acc: 0.9364 - val_loss: 0.1046 - val_acc: 0.9674

Epoch 2/20

55000/55000 [=====] - 9s 164us/step - loss: 0.0827 - acc: 0.9744 - val_loss: 0.0787 - val_acc: 0.9746

Epoch 3/20

55000/55000 [=====] - 9s 163us/step - loss: 0.0573 - acc: 0.9817 - val_loss: 0.0759 - val_acc: 0.9772

Epoch 4/20

55000/55000 [=====] - 9s 164us/step - loss: 0.0411 -
acc: 0.9866 - val_loss: 0.0795 - val_acc: 0.9794
Epoch 5/20
55000/55000 [=====] - 9s 165us/step - loss: 0.0344 -
acc: 0.9881 - val_loss: 0.0816 - val_acc: 0.9756
Epoch 6/20
55000/55000 [=====] - 9s 164us/step - loss: 0.0270 -
acc: 0.9910 - val_loss: 0.0770 - val_acc: 0.9798
Epoch 7/20
55000/55000 [=====] - 9s 167us/step - loss: 0.0228 -
acc: 0.9927 - val_loss: 0.0954 - val_acc: 0.9796
Epoch 8/20
55000/55000 [=====] - 9s 165us/step - loss: 0.0253 -
acc: 0.9918 - val_loss: 0.0875 - val_acc: 0.9784
Epoch 9/20
55000/55000 [=====] - 9s 164us/step - loss: 0.0213 -
acc: 0.9933 - val_loss: 0.0779 - val_acc: 0.9820
Epoch 10/20
55000/55000 [=====] - 9s 164us/step - loss: 0.0150 -
acc: 0.9954 - val_loss: 0.0848 - val_acc: 0.9812
Epoch 11/20
55000/55000 [=====] - 9s 165us/step - loss: 0.0191 -
acc: 0.9940 - val_loss: 0.0892 - val_acc: 0.9798
Epoch 12/20
55000/55000 [=====] - 9s 165us/step - loss: 0.0126 -
acc: 0.9958 - val_loss: 0.0962 - val_acc: 0.9804
Epoch 13/20
55000/55000 [=====] - 9s 166us/step - loss: 0.0181 -
acc: 0.9948 - val_loss: 0.0839 - val_acc: 0.9812
Epoch 14/20
55000/55000 [=====] - 9s 166us/step - loss: 0.0137 -
acc: 0.9957 - val_loss: 0.1121 - val_acc: 0.9790
Epoch 15/20
55000/55000 [=====] - 9s 165us/step - loss: 0.0118 -
acc: 0.9963 - val_loss: 0.0789 - val_acc: 0.9840
Epoch 16/20
55000/55000 [=====] - 9s 165us/step - loss: 0.0091 -
acc: 0.9973 - val_loss: 0.0987 - val_acc: 0.9812
Epoch 17/20
55000/55000 [=====] - 9s 167us/step - loss: 0.0101 -
acc: 0.9971 - val_loss: 0.0824 - val_acc: 0.9826
Epoch 18/20
55000/55000 [=====] - 9s 166us/step - loss: 0.0119 -
acc: 0.9964 - val_loss: 0.1023 - val_acc: 0.9814
Epoch 19/20
55000/55000 [=====] - 9s 166us/step - loss: 0.0141 -
acc: 0.9961 - val_loss: 0.0763 - val_acc: 0.9834
Epoch 20/20


55000/55000 [=====] - 9s 168us/step - loss: 0.0074 -
acc: 0.9979 - val_loss: 0.0964 - val_acc: 0.9830


```
[10]: ## Visualizamos curvas de entrenamiento

plt.figure()
plt.title("Accuracy vs epochs")
plt.plot(history_3h1.history['acc'])
plt.plot(history_3h1.history['val_acc'])
plt.legend(['Training Accuracy', 'Validation Accuracy'])
plt.xlabel("Epoch")
plt.ylabel("Accuracy")

plt.figure()
plt.title("Loss vs epochs")
plt.plot(history_3h1.history['loss'])
plt.plot(history_3h1.history['val_loss'])
plt.legend(['Training Loss', 'Validation Loss'])
plt.xlabel("Epoch")
plt.ylabel("Loss")
```

```
[10]: Text(0,0.5,'Loss')
```





[]:

Red Fully Connected - 5 Capas Ocultas

```
[11]: model_5hl = Sequential()
model_5hl.add(Dense(512, input_shape=(784,), activation="relu"))
model_5hl.add(Dense(1024, activation="relu"))
model_5hl.add(Dense(1024, activation="relu"))
model_5hl.add(Dense(1024, activation="relu"))
model_5hl.add(Dense(512, activation="relu"))
model_5hl.add(Dense(10, activation="softmax"))

model_5hl.compile(
    loss='categorical_crossentropy',
    metrics=['accuracy'],
    optimizer='adam'
)
```

```
[12]: BATCH_SIZE = 128
NUM_EPOCHS = 20

history_5hl = model_5hl.fit(
    X_train, y_train,
    batch_size=BATCH_SIZE,
```



```
epochs=NUM_EPOCHS,  
verbose=1,  
validation_data=(X_val, y_val)  
)
```

Train on 55000 samples, validate on 5000 samples

Epoch 1/20

55000/55000 [=====] - 23s 412us/step - loss: 0.2396 -
acc: 0.9275 - val_loss: 0.1137 - val_acc: 0.9662

Epoch 2/20

55000/55000 [=====] - 22s 397us/step - loss: 0.1053 -
acc: 0.9692 - val_loss: 0.1000 - val_acc: 0.9686

Epoch 3/20

55000/55000 [=====] - 22s 403us/step - loss: 0.0765 -
acc: 0.9772 - val_loss: 0.0919 - val_acc: 0.9752

Epoch 4/20

55000/55000 [=====] - 22s 399us/step - loss: 0.0528 -
acc: 0.9838 - val_loss: 0.0956 - val_acc: 0.9764

Epoch 5/20

55000/55000 [=====] - 22s 399us/step - loss: 0.0497 -
acc: 0.9858 - val_loss: 0.1267 - val_acc: 0.9664

Epoch 6/20

55000/55000 [=====] - 22s 401us/step - loss: 0.0396 -
acc: 0.9885 - val_loss: 0.0741 - val_acc: 0.9822

Epoch 7/20

55000/55000 [=====] - 22s 403us/step - loss: 0.0406 -
acc: 0.9891 - val_loss: 0.1016 - val_acc: 0.9748

Epoch 8/20

55000/55000 [=====] - 22s 400us/step - loss: 0.0312 -
acc: 0.9913 - val_loss: 0.0946 - val_acc: 0.9788

Epoch 9/20

55000/55000 [=====] - 22s 400us/step - loss: 0.0278 -
acc: 0.9920 - val_loss: 0.1148 - val_acc: 0.9768

Epoch 10/20

55000/55000 [=====] - 22s 395us/step - loss: 0.0338 -
acc: 0.9909 - val_loss: 0.0974 - val_acc: 0.9798

Epoch 11/20

55000/55000 [=====] - 22s 395us/step - loss: 0.0183 -
acc: 0.9952 - val_loss: 0.0880 - val_acc: 0.9804

Epoch 12/20

55000/55000 [=====] - 22s 397us/step - loss: 0.0249 -
acc: 0.9936 - val_loss: 0.0967 - val_acc: 0.9804

Epoch 13/20

55000/55000 [=====] - 22s 401us/step - loss: 0.0257 -
acc: 0.9930 - val_loss: 0.0989 - val_acc: 0.9802

Epoch 14/20

55000/55000 [=====] - 22s 401us/step - loss: 0.0224 -

```

acc: 0.9942 - val_loss: 0.1058 - val_acc: 0.9798
Epoch 15/20
55000/55000 [=====] - 22s 395us/step - loss: 0.0215 -
acc: 0.9949 - val_loss: 0.0993 - val_acc: 0.9820
Epoch 16/20
55000/55000 [=====] - 22s 392us/step - loss: 0.0195 -
acc: 0.9950 - val_loss: 0.1253 - val_acc: 0.9784
Epoch 17/20
55000/55000 [=====] - 22s 394us/step - loss: 0.0167 -
acc: 0.9958 - val_loss: 0.1127 - val_acc: 0.9802
Epoch 18/20
55000/55000 [=====] - 22s 392us/step - loss: 0.0165 -
acc: 0.9960 - val_loss: 0.1147 - val_acc: 0.9806
Epoch 19/20
55000/55000 [=====] - 22s 392us/step - loss: 0.0159 -
acc: 0.9961 - val_loss: 0.1269 - val_acc: 0.9804
Epoch 20/20
55000/55000 [=====] - 22s 409us/step - loss: 0.0190 -
acc: 0.9958 - val_loss: 0.1492 - val_acc: 0.9764

```

```

[13]: ## Visualizamos curvas de entrenamiento

plt.figure()
plt.title("Accuracy vs epochs")
plt.plot(history_5hl.history['acc'])
plt.plot(history_5hl.history['val_acc'])
plt.legend(['Training Accuracy', 'Validation Accuracy'])
plt.xlabel("Epoch")
plt.ylabel("Accuracy")


plt.figure()
plt.title("Loss vs epochs")
plt.plot(history_5hl.history['loss'])
plt.plot(history_5hl.history['val_loss'])
plt.legend(['Training Loss', 'Validation Loss'])
plt.xlabel("Epoch")
plt.ylabel("Loss")

```

```

[13]: Text(0,0.5,'Loss')

```



Conclusión: Las dos redes overfittean, pero la que menos overfitting presenta es la red con 3 capas ocultas.

Ejercicio 3

Use la comparación anterior para elegir un modelo.

Realice las predicciones en el conjunto de test y grafique 9 imágenes mal clasificada por el modelo elegido.

```
[14]: score_3hl = model_3hl.evaluate(X_test, y_test, verbose=1)
      print(f'Test Loss      : {score_3hl[0]}')
      print(f'Test Accuracy: {score_3hl[1]}')
```

```
10000/10000 [=====] - 1s 67us/step
Test Loss      : 0.1029824515523373
Test Accuracy: 0.9811
```

```
[15]: y_pred = model_3hl.predict_classes(X_test)

display(y_test)
display(y_pred)
```

```
array([[0., 0., 0., ..., 1., 0., 0.],
       [0., 0., 1., ..., 0., 0., 0.],
       [0., 1., 0., ..., 0., 0., 0.],
       ...,
       [0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.]])
```

```
array([7, 2, 1, ..., 4, 5, 6])
```

```
[16]: enc = OneHotEncoder()
      y_pred_onehot = enc.fit_transform(y_pred.reshape(-1, 1)).toarray()

      display(y_pred_onehot)
```

```
array([[0., 0., 0., ..., 1., 0., 0.],
       [0., 0., 1., ..., 0., 0., 0.],
       [0., 1., 0., ..., 0., 0., 0.],
       ...,
       [0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.]])
```

```

[17]: idxs_wrong_pred = np.array([i for i in range(len(y_test)) if (y_pred_onehot[i] !
    ↪= y_test[i]).any()])


# Visualizamos algunas imagenes
fig, ax = plt.subplots(3, 3, figsize=[15, 15])

for j in range(3):
    for i in range(3):
        for idx in np.random.choice(idxs_wrong_pred, 9):
            image = X_test[idx]
            ground_truth = int(enc.inverse_transform(y_test[idx].reshape(1,
    ↪-1))[0][0])
            pred = int(enc.inverse_transform(y_pred_onehot[idx].reshape(1,
    ↪-1))[0][0])

            ax[i,j].imshow(image.reshape(28, 28), 'gray')
            ax[i,j].set_title(f'Pred: {pred} / GT: {ground_truth}', color='r')
            ax[i,j].set_xticks([])
            ax[i,j].set_yticks([])

# fig.subplots_adjust(wspace=0.1, hspace=0)
# plt.title("9 imágenes mal clasificada en el conjunto de test")
plt.show()

```



Ejercicio 4

Defina dos redes convolucionales diferentes que sean capaces de clasificar la base MNIST.

Utilice para entrenar ambas la base de entrenamiento aumentada (por $21=3 \times 3 \times 3$) utilizando rotación, escala y traslación.

```
[18]: X_train = mnist.train.images
      y_train = mnist.train.labels

      X_val = mnist.validation.images
      y_val = mnist.validation.labels
```

```
X_test = mnist.test.images
y_test = mnist.test.labels
```

```
[19]: X_train = X_train.reshape(55000, 28, 28, 1)
      X_val = X_val.reshape(5000, 28, 28, 1)
      X_test = X_test.reshape(10000, 28, 28, 1)
```

```
[20]: print(f"X_train shape: {X_train.shape}")
      print(f"y_train shape: {y_train.shape}")
      print()
      print(f"X_val shape : {X_val.shape}")
      print(f"y_val shape : {y_val.shape}")
      print()
      print(f"X_test shape : {X_test.shape}")
      print(f"y_test shape : {y_test.shape}")
```

```
X_train shape: (55000, 28, 28, 1)
y_train shape: (55000, 10)
```

```
X_val shape : (5000, 28, 28, 1)
y_val shape : (5000, 10)
```

```
X_test shape : (10000, 28, 28, 1)
y_test shape : (10000, 10)
```

```
[21]: ##### Some Constants #####
      IMG_ROWS = 28
      IMG_COLS = 28
      NUM_CLASSES = 10

      BATCH_SIZE = 128
      NUM_EPOCHS = 20
```

```
[22]: ### Data Augmentation ###

      # Training data generator
      datagen_train = ImageDataGenerator(
          # rescale=1./255, # We also can make a rescale on the data
          rotation_range=10,
          width_shift_range=0.2,
          height_shift_range=0.2
      )

      # Validation data generator
      datagen_val = ImageDataGenerator(
          # rescale=1./255, # We also can make a rescale on the data
      )
```

```
#####
```

CNN sin dropout

```
[23]: inputs = Input(shape=(IMG_ROWS, IMG_COLS, 1))

conv1 = Conv2D(64, 5, activation="relu")(inputs)
maxPool1 = MaxPooling2D(pool_size=(2, 2))(conv1)

conv2 = Conv2D(128, 3, activation="relu")(maxPool1)
maxPool2 = MaxPooling2D(pool_size=(2, 2))(conv2)

flatten = Flatten()(maxPool2)

dense1 = Dense(256, activation="sigmoid")(flatten)
dense2 = Dense(64, activation="sigmoid")(dense1)

outputs = Dense(NUM_CLASSES, activation="softmax")(dense2)

model_cnn1 = Model(
    inputs=inputs,
    outputs=outputs
)

model_cnn1.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy']
)
```

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:4267: The name tf.nn.max_pool is deprecated. Please use tf.nn.max_pool2d instead.

```
[24]: model_cnn1.summary()
```

Model: "model_1"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 28, 28, 1)	0
conv2d_1 (Conv2D)	(None, 24, 24, 64)	1664
max_pooling2d_1 (MaxPooling2D)	(None, 12, 12, 64)	0

conv2d_2 (Conv2D)	(None, 10, 10, 128)	73856

max_pooling2d_2 (MaxPooling2)	(None, 5, 5, 128)	0

flatten_1 (Flatten)	(None, 3200)	0

dense_11 (Dense)	(None, 256)	819456

dense_12 (Dense)	(None, 64)	16448

dense_13 (Dense)	(None, 10)	650
=====		
Total params: 912,074		
Trainable params: 912,074		
Non-trainable params: 0		

```
[25]: history_cnn1 = model_cnn1.fit_generator(
    datagen_train.flow(X_train, y_train, batch_size=BATCH_SIZE),
    steps_per_epoch=X_train.shape[0] // BATCH_SIZE,
    epochs=NUM_EPOCHS,
    verbose=1,
    validation_data=datagen_val.flow(X_val, y_val, batch_size=BATCH_SIZE),
    validation_steps=X_val.shape[0] // BATCH_SIZE
)
```

```
Epoch 1/20
429/429 [=====] - 52s 122ms/step - loss: 0.7443 - acc:
0.7946 - val_loss: 0.1806 - val_acc: 0.9529
Epoch 2/20
429/429 [=====] - 52s 121ms/step - loss: 0.1993 - acc:
0.9455 - val_loss: 0.0984 - val_acc: 0.9733
Epoch 3/20
429/429 [=====] - 52s 121ms/step - loss: 0.1323 - acc:
0.9617 - val_loss: 0.0674 - val_acc: 0.9803
Epoch 4/20
429/429 [=====] - 52s 120ms/step - loss: 0.1043 - acc:
0.9689 - val_loss: 0.0541 - val_acc: 0.9854
Epoch 5/20
429/429 [=====] - 52s 121ms/step - loss: 0.0828 - acc:
0.9750 - val_loss: 0.0444 - val_acc: 0.9869
Epoch 6/20
429/429 [=====] - 52s 121ms/step - loss: 0.0711 - acc:
0.9787 - val_loss: 0.0542 - val_acc: 0.9844
Epoch 7/20
429/429 [=====] - 52s 120ms/step - loss: 0.0651 - acc:
0.9800 - val_loss: 0.0420 - val_acc: 0.9865
Epoch 8/20
```

```

429/429 [=====] - 51s 120ms/step - loss: 0.0574 - acc:
0.9826 - val_loss: 0.0391 - val_acc: 0.9856
Epoch 9/20
429/429 [=====] - 51s 120ms/step - loss: 0.0547 - acc:
0.9834 - val_loss: 0.0578 - val_acc: 0.9840
Epoch 10/20
429/429 [=====] - 51s 120ms/step - loss: 0.0470 - acc:
0.9855 - val_loss: 0.0412 - val_acc: 0.9897
Epoch 11/20
429/429 [=====] - 51s 120ms/step - loss: 0.0484 - acc:
0.9850 - val_loss: 0.0445 - val_acc: 0.9852
Epoch 12/20
429/429 [=====] - 51s 120ms/step - loss: 0.0452 - acc:
0.9861 - val_loss: 0.0440 - val_acc: 0.9865
Epoch 13/20
429/429 [=====] - 51s 120ms/step - loss: 0.0453 - acc:
0.9859 - val_loss: 0.0341 - val_acc: 0.9889
Epoch 14/20
429/429 [=====] - 51s 120ms/step - loss: 0.0430 - acc:
0.9869 - val_loss: 0.0245 - val_acc: 0.9922
Epoch 15/20
429/429 [=====] - 52s 121ms/step - loss: 0.0383 - acc:
0.9885 - val_loss: 0.0421 - val_acc: 0.9865
Epoch 16/20
429/429 [=====] - 51s 120ms/step - loss: 0.0390 - acc:
0.9878 - val_loss: 0.0323 - val_acc: 0.9906
Epoch 17/20
429/429 [=====] - 52s 120ms/step - loss: 0.0380 - acc:
0.9882 - val_loss: 0.0289 - val_acc: 0.9899
Epoch 18/20
429/429 [=====] - 52s 121ms/step - loss: 0.0340 - acc:
0.9895 - val_loss: 0.0257 - val_acc: 0.9920
Epoch 19/20
429/429 [=====] - 52s 121ms/step - loss: 0.0338 - acc:
0.9892 - val_loss: 0.0261 - val_acc: 0.9908
Epoch 20/20
429/429 [=====] - 52s 121ms/step - loss: 0.0329 - acc:
0.9896 - val_loss: 0.0294 - val_acc: 0.9926

```

[26]: *## Visualizamos curvas de entrenamiento*

```


plt.figure()
plt.title("Accuracy vs epochs")
plt.plot(history_cnn1.history['acc'])
plt.plot(history_cnn1.history['val_acc'])
plt.legend(['Training Accuracy', 'Validation Accuracy'])
plt.xlabel("Epoch")


```

```
plt.ylabel("Accuracy")

plt.figure()
plt.title("Loss vs epochs")
plt.plot(history_cnn1.history['loss'])
plt.plot(history_cnn1.history['val_loss'])
plt.legend(['Training Loss', 'Validation Loss'])
plt.xlabel("Epoch")
plt.ylabel("Loss")
```

[26]: Text(0,0.5,'Loss')





```
[30]: score_cnn1 = model_cnn1.evaluate(X_test, y_test, verbose=1)
      print(f'Test Loss      : {score_cnn1[0]}')
      print(f'Test Accuracy: {score_cnn1[1]}')
```

```
10000/10000 [=====] - 3s 266us/step
Test Loss      : 0.028125276191858575
Test Accuracy: 0.9915
```

```
[ ]:
```

CNN con dropout

```
[27]: inputs = Input(shape=(IMG_ROWS, IMG_COLS, 1))

conv1 = Conv2D(64, 5, activation=None)(inputs)
bn1 = BatchNormalization()(conv1)
activation1 = Activation("relu")(bn1)
maxPool1 = MaxPooling2D(pool_size=(2, 2))(activation1)

conv2 = Conv2D(128, 5, activation=None)(maxPool1)
bn2 = BatchNormalization()(conv2)
activation2 = Activation("relu")(bn2)
maxPool2 = MaxPooling2D(pool_size=(2, 2))(activation2)
```

```

flatten = Flatten()(maxPool2)

dense1 = Dense(256, activation="sigmoid")(flatten)
dense1 = Dropout(0.5)(dense1)

dense2 = Dense(64, activation="sigmoid")(dense1)
dense2 = Dropout(0.5)(dense2)

outputs = Dense(NUM_CLASSES, activation="softmax")(dense2)

model_cnn2 = Model(
    inputs=inputs,
    outputs=outputs
)

model_cnn2.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy']
)

```

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:2041: The name tf.nn.fused_batch_norm is deprecated. Please use tf.compat.v1.nn.fused_batch_norm instead.

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:148: The name tf.placeholder_with_default is deprecated. Please use tf.compat.v1.placeholder_with_default instead.

WARNING:tensorflow:From /home/ubuntu/miniconda3/envs/mlip/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:3733: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

```

[28]: history_cnn2 = model_cnn2.fit_generator(
    datagen_train.flow(X_train, y_train, batch_size=BATCH_SIZE),
    steps_per_epoch=X_train.shape[0] // BATCH_SIZE,
    epochs=NUM_EPOCHS,
    validation_data=datagen_val.flow(X_val, y_val, batch_size=BATCH_SIZE),
    validation_steps=X_val.shape[0] // BATCH_SIZE
)

```

Epoch 1/20

429/429 [=====] - 122s 285ms/step - loss: 1.0040 - acc: 0.6955 - val_loss: 0.1691 - val_acc: 0.9663
Epoch 2/20
429/429 [=====] - 121s 283ms/step - loss: 0.2789 - acc: 0.9346 - val_loss: 0.1184 - val_acc: 0.9680
Epoch 3/20
429/429 [=====] - 121s 283ms/step - loss: 0.1983 - acc: 0.9516 - val_loss: 0.1295 - val_acc: 0.9622
Epoch 4/20
429/429 [=====] - 122s 284ms/step - loss: 0.1598 - acc: 0.9584 - val_loss: 0.1322 - val_acc: 0.9602
Epoch 5/20
429/429 [=====] - 122s 284ms/step - loss: 0.1443 - acc: 0.9627 - val_loss: 0.0407 - val_acc: 0.9881
Epoch 6/20
429/429 [=====] - 121s 283ms/step - loss: 0.1271 - acc: 0.9677 - val_loss: 0.0788 - val_acc: 0.9766
Epoch 7/20
429/429 [=====] - 121s 283ms/step - loss: 0.1221 - acc: 0.9684 - val_loss: 0.0918 - val_acc: 0.9709
Epoch 8/20
429/429 [=====] - 121s 283ms/step - loss: 0.1135 - acc: 0.9710 - val_loss: 0.0438 - val_acc: 0.9854
Epoch 9/20
429/429 [=====] - 122s 284ms/step - loss: 0.1016 - acc: 0.9733 - val_loss: 0.0424 - val_acc: 0.9860
Epoch 10/20
429/429 [=====] - 121s 283ms/step - loss: 0.0956 - acc: 0.9745 - val_loss: 0.0957 - val_acc: 0.9684
Epoch 11/20
429/429 [=====] - 121s 282ms/step - loss: 0.0938 - acc: 0.9760 - val_loss: 0.0363 - val_acc: 0.9889
Epoch 12/20
429/429 [=====] - 121s 283ms/step - loss: 0.0883 - acc: 0.9775 - val_loss: 0.0347 - val_acc: 0.9893
Epoch 13/20
429/429 [=====] - 121s 283ms/step - loss: 0.0867 - acc: 0.9775 - val_loss: 0.0796 - val_acc: 0.9770
Epoch 14/20
429/429 [=====] - 121s 283ms/step - loss: 0.0820 - acc: 0.9790 - val_loss: 0.0527 - val_acc: 0.9838
Epoch 15/20
429/429 [=====] - 121s 282ms/step - loss: 0.0811 - acc: 0.9789 - val_loss: 0.0385 - val_acc: 0.9887
Epoch 16/20
429/429 [=====] - 121s 283ms/step - loss: 0.0790 - acc: 0.9794 - val_loss: 0.0320 - val_acc: 0.9901
Epoch 17/20

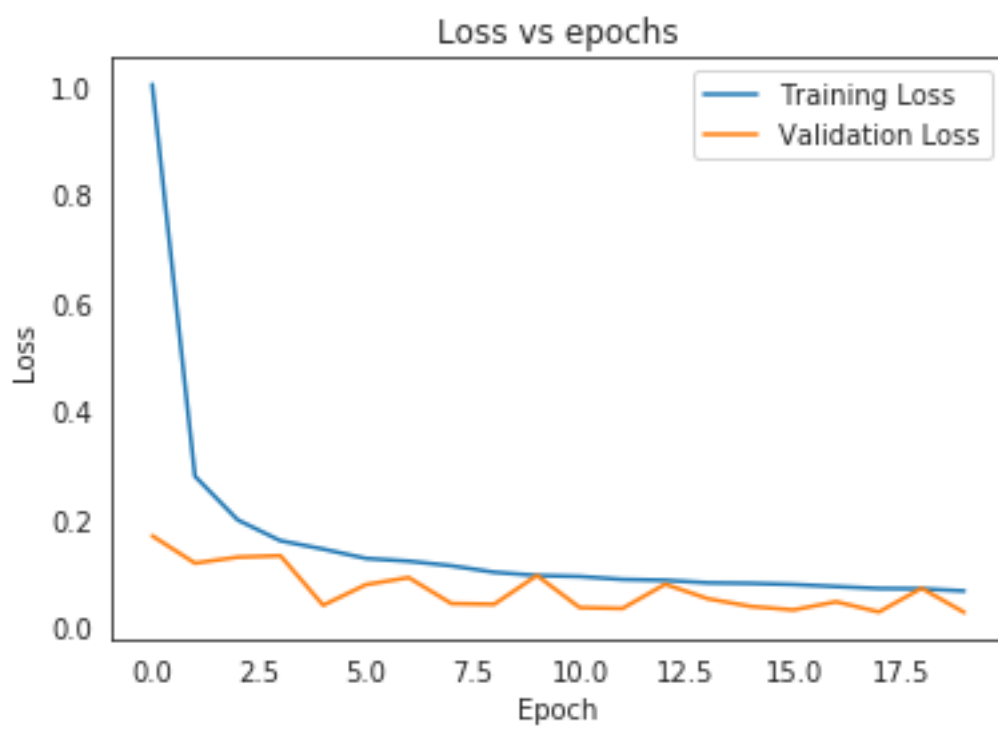

```
429/429 [=====] - 121s 283ms/step - loss: 0.0752 - acc:
0.9808 - val_loss: 0.0471 - val_acc: 0.9838
Epoch 18/20
429/429 [=====] - 121s 283ms/step - loss: 0.0717 - acc:
0.9820 - val_loss: 0.0283 - val_acc: 0.9920
Epoch 19/20
429/429 [=====] - 122s 283ms/step - loss: 0.0711 - acc:
0.9829 - val_loss: 0.0721 - val_acc: 0.9795
Epoch 20/20
429/429 [=====] - 121s 283ms/step - loss: 0.0670 - acc:
0.9829 - val_loss: 0.0273 - val_acc: 0.9914
```

```
[29]: ## Visualizamos curvas de entrenamiento

plt.figure()
plt.title("Accuracy vs epochs")
plt.plot(history_cnn2.history['acc'])
plt.plot(history_cnn2.history['val_acc'])
plt.legend(['Training Accuracy', 'Validation Accuracy'])
plt.xlabel("Epoch")
plt.ylabel("Accuracy")

plt.figure()
plt.title("Loss vs epochs")
plt.plot(history_cnn2.history['loss'])
plt.plot(history_cnn2.history['val_loss'])
plt.legend(['Training Loss', 'Validation Loss'])
plt.xlabel("Epoch")
plt.ylabel("Loss")
```

```
[29]: Text(0,0.5,'Loss')
```




```
[31]: score_cnn2 = model_cnn2.evaluate(X_test, y_test, verbose=1)
      print(f'Test Loss      : {score_cnn2[0]}')
      print(f'Test Accuracy: {score_cnn2[1]}')
```

```
10000/10000 [=====] - 6s 598us/step
Test Loss      : 0.03212649761128705
Test Accuracy: 0.9908
```

Ejercicio 5

Describe en detalle una posible aplicación de las redes convolucionales vistas en el taller a un problema de clasificación en ciencia, el estado o en la industria.

Atributador de Fashion

Una posible aplicación de las redes convolucionales en la industria es un atributador de ropa, en el cual la red sea capaz de aprender si una remera es manga corta o manga largo, si tiene el cuello redondo o en escote V. Además del color de la ropa.

Este uso es muy útil en los ecommerce cuando alguien quiere vender una ropa pero no detalla las características de la misma, siendo la red muy útil en este caso ayudando a los usuarios a buscar prendas por sus atributos.

```
[ ]:
```